

EX 3 IMPLEMENTATION OF CONVOLUTION NEURAL NETWORK (CNN)

AIM:

To implement a Convolutional Neural Network (CNN) using PyTorch for classifying CIFAR-10 images and visualize the learned convolution filters.

PROGRAM:

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# Step 1: Load and preprocess CIFAR-10
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5),
                          (0.5, 0.5, 0.5))
])

trainset = torchvision.datasets.CIFAR10(
    root='./data', train=True, download=True, transform=transform
)
testset = torchvision.datasets.CIFAR10(
    root='./data', train=False, download=True, transform=transform
)

trainloader = torch.utils.data.DataLoader(
    trainset, batch_size=64, shuffle=True
)
testloader = torch.utils.data.DataLoader(
    testset, batch_size=64, shuffle=False
)

# Step 2: Define CNN
class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()

        self.conv1 = nn.Conv2d(3, 16, 3, padding=1)
        self.conv2 = nn.Conv2d(16, 32, 3, padding=1)

        self.pool = nn.MaxPool2d(2, 2)
        self.relu = nn.ReLU()

        self.fc1 = nn.Linear(32 * 8 * 8, 128)
        self.fc2 = nn.Linear(128, 10)
```

```

def forward(self, x):
    x = self.pool(self.relu(self.conv1(x)))
    x = self.pool(self.relu(self.conv2(x)))

    x = x.view(-1, 32 * 8 * 8)

    x = self.relu(self.fc1(x))
    x = self.fc2(x)

    return x

# Step 3: Train the CNN
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model = SimpleCNN().to(device)

lossfn = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

for epoch in range(10):
    running_loss = 0.0

    for inputs, labels in trainloader:
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = model(inputs)
        loss = lossfn(outputs, labels)

        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    print(
        f"Epoch {epoch+1}, "
        f"Loss: {running_loss / len(trainloader):.4f}"
    )

# Step 4: Evaluate the model
correct = 0
total = 0

with torch.no_grad():
    for images, labels in testloader:
        images, labels = images.to(device), labels.to(device)

        outputs = model(images)
        _, predicted = outputs.max(1)

```

```

total += labels.size(o)
correct += (predicted == labels).sum().item()

print(f"Accuracy on test data: {100 * correct / total:.2f}%")

# Step 5: Visualize learned filters
def visualize_filters(layer, n_filters=8):
    filters = layer.weight.data.clone().cpu()

    fig, axs = plt.subplots(
        1, n_filters, figsize=(15, 4)
    )

    for i in range(n_filters):
        f = filters[i]
        f = (f - f.min()) / (f.max() - f.min())

        axs[i].imshow(f.permute(1, 2, 0))
        axs[i].axis('off')
        axs[i].set_title(f'Filter {i+1}')

    plt.tight_layout()
    plt.show()

visualize_filters(model.conv1)


```

OUTPUT:

```

Epoch 1, Loss: 1.4516
Epoch 2, Loss: 1.1065
Epoch 3, Loss: 0.9556
Epoch 4, Loss: 0.8525
Epoch 5, Loss: 0.7670
Epoch 6, Loss: 0.6954
Epoch 7, Loss: 0.6297
Epoch 8, Loss: 0.5653
Epoch 9, Loss: 0.5043
Epoch 10, Loss: 0.4522
Accuracy on test data: 69.30%

```



RESULT:

The CNN was successfully trained on the CIFAR-10 dataset, achieved good classification accuracy, and the learned convolution filters were visualized successfully.

